

Titanic Dataset Preprocessing Pipeline

This script preprocesses the Titanic dataset by:

1. Handling missing values
2. Removing outliers
3. Encoding categorical variables
4. Scaling numerical features

Output: Clean dataset ready for modeling

Step:1 Import libraries

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import OneHotEncoder
```

Step:2 Overview of dataset

```
# Load the dataset
titanic_df = sns.load_dataset('titanic')
titanic_df
```

	survived	pclass	sex	age	sibsp	parch	fare	embarked
class \								
0	0	3	male	22.0	1	0	7.2500	S
Third								
1	1	1	female	38.0	1	0	71.2833	C
First								
2	1	3	female	26.0	0	0	7.9250	S
Third								
3	1	1	female	35.0	1	0	53.1000	S
First								
4	0	3	male	35.0	0	0	8.0500	S
Third								
..
...								
886	0	2	male	27.0	0	0	13.0000	S
Second								
887	1	1	female	19.0	0	0	30.0000	S
First								
888	0	3	female	NaN	1	2	23.4500	S
Third								
889	1	1	male	26.0	0	0	30.0000	C
First								
890	0	3	male	32.0	0	0	7.7500	Q

	who	adult_male	deck	embark_town	alive	alone
0	man	True	NaN	Southampton	no	False
1	woman	False	C	Cherbourg	yes	False
2	woman	False	NaN	Southampton	yes	True
3	woman	False	C	Southampton	yes	False
4	man	True	NaN	Southampton	no	True
...
886	man	True	NaN	Southampton	no	True
887	woman	False	B	Southampton	yes	True
888	woman	False	NaN	Southampton	no	False
889	man	True	C	Cherbourg	yes	True
890	man	True	NaN	Queenstown	no	True

Step:3 Understanding the Data

```
Rows in the dataset: [891]
Columns in the dataset: [15]
```

survived	pclass	sex	age	sibsp	parch	fare	embarked
0	3	male	22.0	1	0	7.2500	S
1	1	female	38.0	1	0	71.2833	C
2	3	female	26.0	0	0	7.9250	S
3	1	female	35.0	1	0	53.1000	S
4	3	male	35.0	0	0	8.0500	S

	who	adult_male	deck	embark_town	alive	alone
0	man	True	NaN	Southampton	no	False
1	woman	False	C	Cherbourg	yes	False
2	woman	False	NaN	Southampton	yes	True
3	woman	False	C	Southampton	yes	False
4	man	True	NaN	Southampton	no	True

```
survived  pclass    sex   age  sibsp  parch   fare embarked
class \
```

886	0	2	male	27.0	0	0	13.00	S
Second								
887	1	1	female	19.0	0	0	30.00	S
First								
888	0	3	female	NaN	1	2	23.45	S
Third								
889	1	1	male	26.0	0	0	30.00	C
First								
890	0	3	male	32.0	0	0	7.75	Q
Third								

	who	adult_male	deck	embark_town	alive	alone
886	man	True	NaN	Southampton	no	True
887	woman	False	B	Southampton	yes	True
888	woman	False	NaN	Southampton	no	False
889	man	True	C	Cherbourg	yes	True
890	man	True	NaN	Queenstown	no	True

```
titanic_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 891 entries, 0 to 890
```

```
Data columns (total 15 columns):
```

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	survived	891 non-null	int64
1	pclass	891 non-null	int64
2	sex	891 non-null	object
3	age	714 non-null	float64
4	sibsp	891 non-null	int64
5	parch	891 non-null	int64
6	fare	891 non-null	float64
7	embarked	889 non-null	object
8	class	891 non-null	category
9	who	891 non-null	object
10	adult_male	891 non-null	bool
11	deck	203 non-null	category
12	embark_town	889 non-null	object
13	alive	891 non-null	object
14	alone	891 non-null	bool

```
dtypes: bool(2), category(2), float64(2), int64(4), object(5)
```

```
memory usage: 80.7+ KB
```

```
titanic_df.describe()
```

	survived	pclass	age	sibsp	parch
fare					
count	891.000000	891.000000	714.000000	891.000000	891.000000
891.000000					
mean	0.383838	2.308642	29.699118	0.523008	0.381594

32.204208					
std	0.486592	0.836071	14.526497	1.102743	0.806057
49.693429					
min	0.000000	1.000000	0.420000	0.000000	0.000000
0.000000					
25%	0.000000	2.000000	20.125000	0.000000	0.000000
7.910400					
50%	0.000000	3.000000	28.000000	0.000000	0.000000
14.454200					
75%	1.000000	3.000000	38.000000	1.000000	0.000000
31.000000					
max	1.000000	3.000000	80.000000	8.000000	6.000000
512.329200					

Step4:Handling missing values

```
# Check for missing values
titanic_df.isnull().sum()

survived      0
pclass        0
sex           0
age          177
sibsp         0
parch         0
fare          0
embarked      2
class         0
who           0
adult_male    0
deck         688
embark_town   2
alive         0
alone         0
dtype: int64

# missing value percentage
print(titanic_df.isnull().sum()/len(titanic_df)*100)

survived      0.000000
pclass        0.000000
sex           0.000000
age          19.865320
sibsp         0.000000
parch         0.000000
fare          0.000000
embarked      0.224467
class         0.000000
who           0.000000
adult_male    0.000000
```

```
deck          77.216611
embark_town    0.224467
alive         0.000000
alone         0.000000
dtype: float64
```

Missing values in columns like Age, Deck, Embarked and Embark_town

```
titanic_df['age'].fillna(titanic_df['age'].median(),inplace=True)
```

```
titanic_df.drop(columns=['deck'],inplace=True)
```

```
titanic_df['embarked'].fillna(titanic_df['embarked'].mode()[0],inplace=True)
```

```
titanic_df['embark_town'].fillna(titanic_df['embark_town'].mode()[0],inplace=True)
```

C:\Users\sanrkin\AppData\Local\Temp\ipykernel_9344\577215265.py:1:
FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
titanic_df['age'].fillna(titanic_df['age'].median(),inplace=True)
```

C:\Users\sanrkin\AppData\Local\Temp\ipykernel_9344\577215265.py:5:
FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
titanic_df['embarked'].fillna(titanic_df['embarked'].mode()[0],inplace=True)
```

C:\Users\sanrkin\AppData\Local\Temp\ipykernel_9344\577215265.py:7:
FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
titanic_df['embark_town'].fillna(titanic_df['embark_town'].mode()[0],inplace=True)
```

```
titanic_df.isnull().sum()
```

```
survived      0
pclass        0
sex            0
age           0
sibsp         0
parch         0
fare          0
embarked      0
class         0
who           0
adult_male    0
embark_town   0
alive         0
alone        0
dtype: int64
```

```
# Drop irrelevant columns
```

```
titanic_df.drop(columns=['parch','sibsp'], inplace=True)
```

```
# Distribution of numerical features
```


```
plt.figure(figsize=(10, 8))
sns.histplot(data=titanic_df, x='age', kde=True, color='blue',
label='Age')
plt.grid(True, linestyle='--', alpha=0.5)
plt.legend()
plt.show()
```

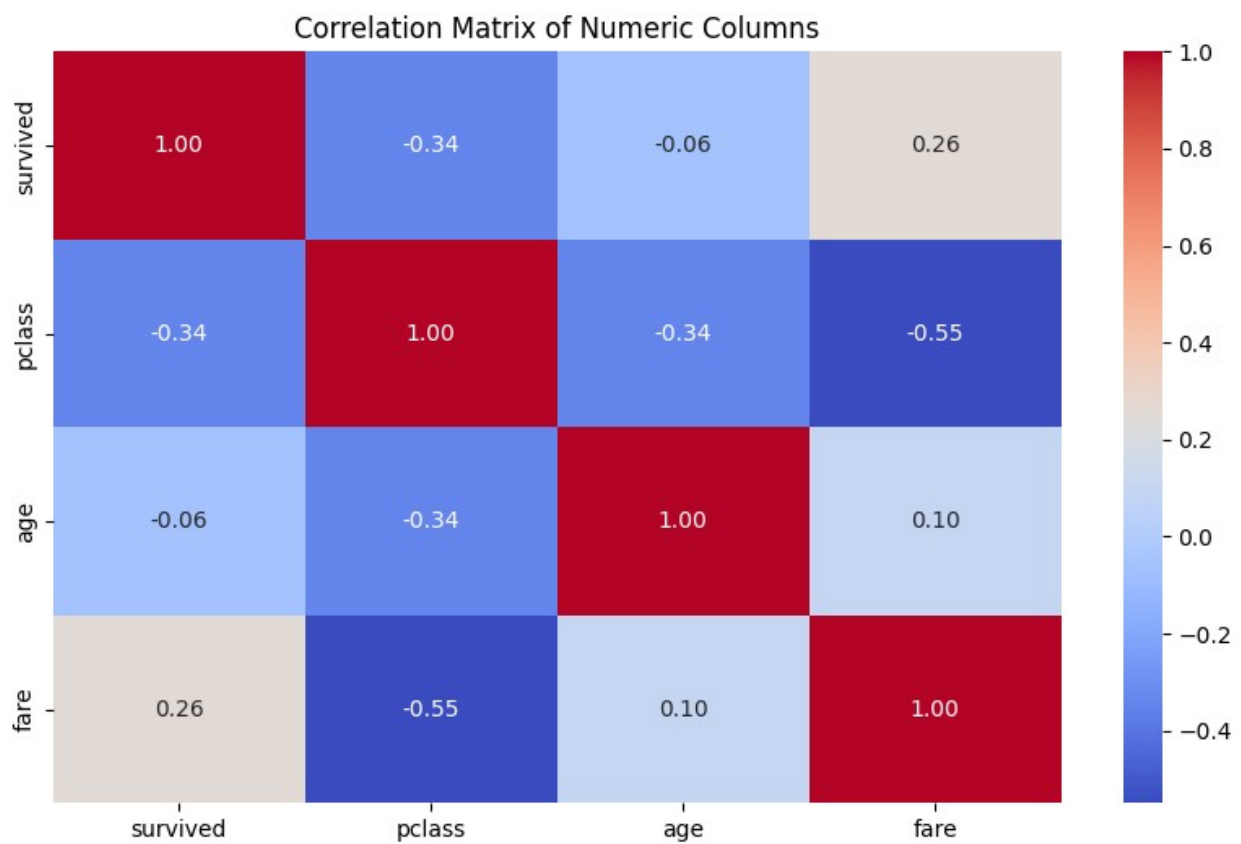

```
# Plot histogram for 'fare'
```

```
plt.figure(figsize=(10, 6))
sns.histplot(data=titanic_df, x='fare', kde=True, color='orange',
label='Fare')
plt.grid(True, linestyle='--', alpha=0.5)
plt.legend()
plt.show()
```

```
numeric_df = titanic_df.select_dtypes(include=['number'])

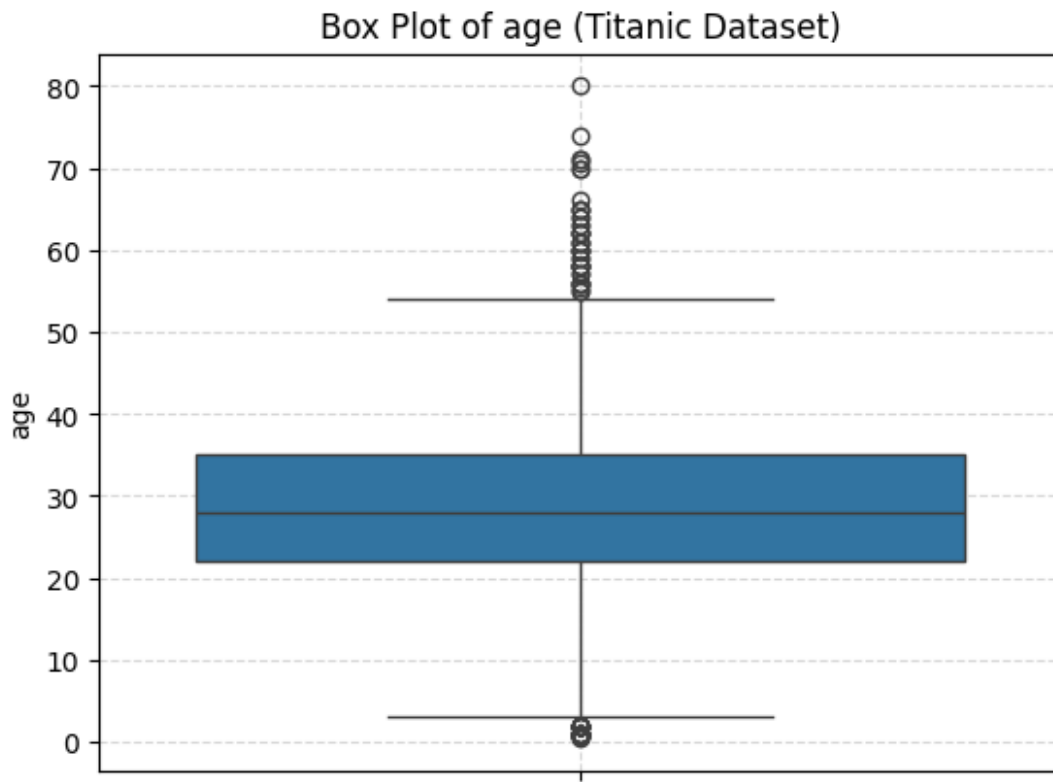
# Plot the correlation matrix
plt.figure(figsize=(10, 6))
sns.heatmap(numeric_df.corr(), annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation Matrix of Numeric Columns')
plt.show()
```

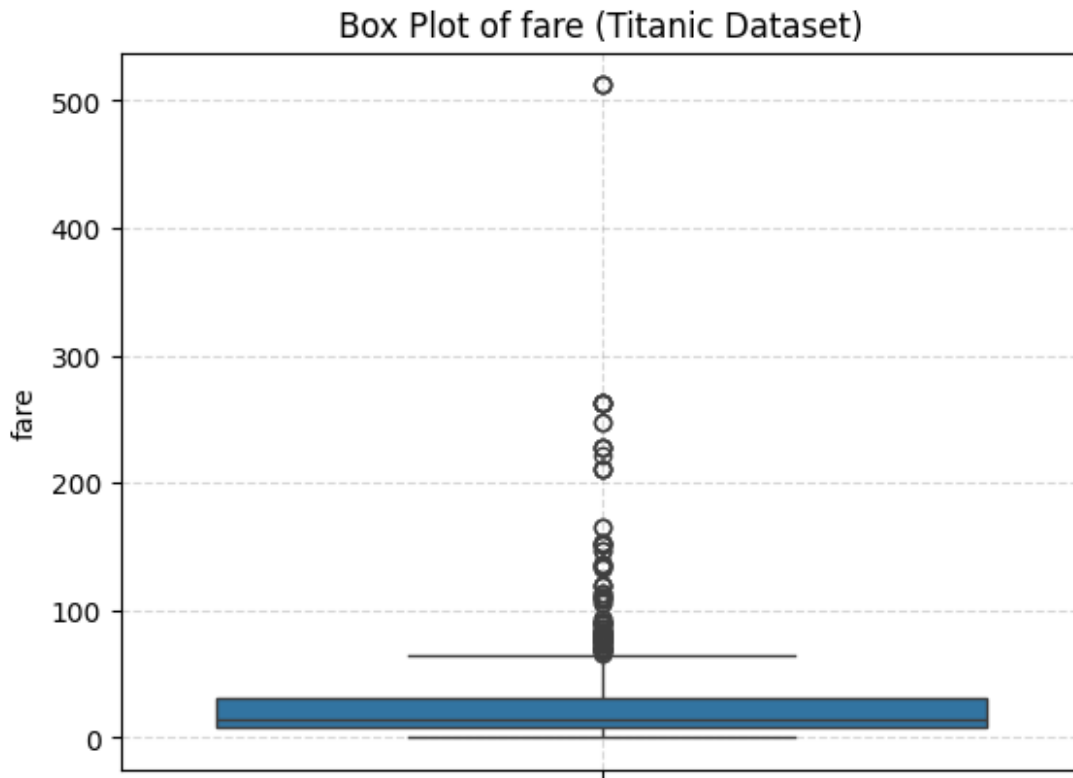




Step5:Handling outliers

```
numerical_columns = ['age', 'fare']  
  
for col in numerical_columns:  
    sns.boxplot(data=titanic_df, y=col)  
    plt.title(f"Box Plot of {col} (Titanic Dataset)")  
    plt.grid(True, linestyle='--', alpha=0.5)  
    plt.show()
```





we can see there are outliers in the age and fare columns

```
def remove_outliers(df, column):
    Q1 = df[column].quantile(0.25)
    Q3 = df[column].quantile(0.75)
    IQR = Q3 - Q1


    # Define the bounds for filtering
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR

    # Filter the dataset to keep only non-outlier data
    df_filtered = df[(df[column] >= lower_bound) & (df[column] <=
upper_bound)]
    return df_filtered


titanic_data_cleaned = remove_outliers(titanic_df, 'age')
titanic_data_cleaned = remove_outliers(titanic_data_cleaned, 'fare')

for col in numerical_columns:
    sns.boxplot(data=titanic_data_cleaned, y=col)
    plt.title(f"Box Plot of '{col}' after removing outliers ")
    plt.grid(True, linestyle='--', alpha=0.5)
    plt.show()
```

Box Plot of 'age' after removing outliers

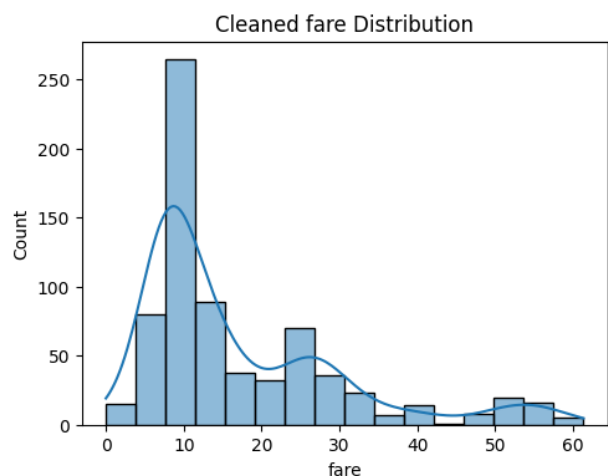

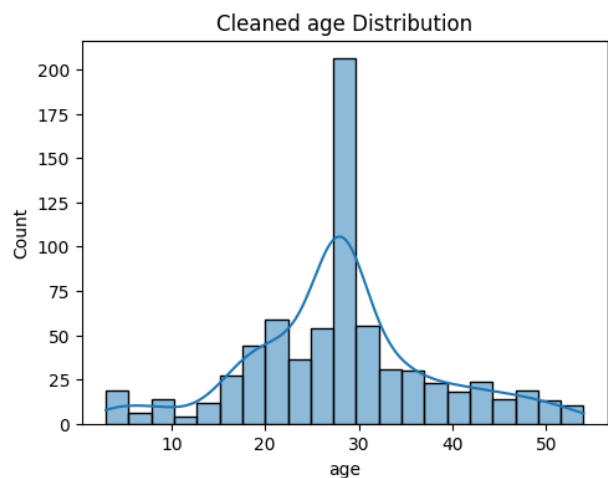



Box Plot of 'fare' after removing outliers



```
def plot_distribution_comparison(df_original, df_cleaned, column):
    plt.figure(figsize=(12, 4))
    plt.subplot(1, 2, 1)
    sns.histplot(df_original[column], kde=True)
    plt.title(f'Original {column} Distribution')
    plt.subplot(1, 2, 2)
    sns.histplot(df_cleaned[column], kde=True)
    plt.title(f'Cleaned {column} Distribution')
    plt.show()

# Use for each numerical column
for col in numerical_columns:
    plot_distribution_comparison(titanic_df, titanic_data_cleaned,
                               col)
```



Step 6: Dealing with Categorical Data

titanic_data_cleaned

who \	survived	pclass	sex	age	fare	embarked	class	
0	0	3	male	22.0	7.2500	S	Third	man
2	1	3	female	26.0	7.9250	S	Third	woman
3	1	1	female	35.0	53.1000	S	First	woman
4	0	3	male	35.0	8.0500	S	Third	man
5	0	3	male	28.0	8.4583	Q	Third	man
..
886	0	2	male	27.0	13.0000	S	Second	man
887	1	1	female	19.0	30.0000	S	First	woman
888	0	3	female	28.0	23.4500	S	Third	woman
889	1	1	male	26.0	30.0000	C	First	man
890	0	3	male	32.0	7.7500	Q	Third	man

	adult_male	embark_town	alive	alone
0	True	Southampton	no	False
2	False	Southampton	yes	True
3	False	Southampton	yes	False
4	True	Southampton	no	True
5	True	Queenstown	no	True
..
886	True	Southampton	no	True
887	False	Southampton	yes	True
888	False	Southampton	no	False
889	True	Cherbourg	yes	True
890	True	Queenstown	no	True

[718 rows x 12 columns]

Create separate encoders

le_sex = LabelEncoder()

le_who = LabelEncoder()

Encode sex

titanic_data_cleaned['sex_encoded'] =

le_sex.fit_transform(titanic_data_cleaned['sex'])

Encode who

titanic_data_cleaned['who_encoded'] =

le_who.fit_transform(titanic_data_cleaned['who'])

```
# If you need to know the mapping:
print("Sex categories:", le_sex.classes_)
print("Who categories:", le_who.classes_)
```

```
Sex categories: ['female' 'male']
Who categories: ['child' 'man' 'woman']
```

Step 7: Convert Categorical Columns into Numerical

```
# Initialize OneHotEncoder
ohe = OneHotEncoder(sparse_output=False) # Use sparse=False for a
dense array

# Fit and transform the 'class' column
encoded_class = ohe.fit_transform(titanic_data_cleaned[['class']])
#2d

# Convert the encoded array to a DataFrame with proper column names
class_encoded_df = pd.DataFrame(encoded_class,
columns=ohe.get_feature_names_out(['class']))

# Ensure the encoded DataFrame has the same index as the original
DataFrame
class_encoded_df.index = titanic_data_cleaned.index

# Concatenate the encoded DataFrame with the original DataFrame
df_titanic = pd.concat([titanic_data_cleaned, class_encoded_df],
axis=1)

# Display the resulting DataFrame
df_titanic
```

	survived	pclass	sex	age	fare	embarked	class	
who \								
0	0	3	male	22.0	7.2500	S	Third	man
2	1	3	female	26.0	7.9250	S	Third	woman
3	1	1	female	35.0	53.1000	S	First	woman
4	0	3	male	35.0	8.0500	S	Third	man
5	0	3	male	28.0	8.4583	Q	Third	man
...
886	0	2	male	27.0	13.0000	S	Second	man

887	1	1	female	19.0	30.0000	S	First	woman
888	0	3	female	28.0	23.4500	S	Third	woman
889	1	1	male	26.0	30.0000	C	First	man
890	0	3	male	32.0	7.7500	Q	Third	man

	adult_male	embark_town	alive	alone	sex_encoded	who_encoded	\
0	True	Southampton	no	False	1	1	
2	False	Southampton	yes	True	0	2	
3	False	Southampton	yes	False	0	2	
4	True	Southampton	no	True	1	1	
5	True	Queenstown	no	True	1	1	
..	
886	True	Southampton	no	True	1	1	
887	False	Southampton	yes	True	0	2	
888	False	Southampton	no	False	0	2	
889	True	Cherbourg	yes	True	1	1	
890	True	Queenstown	no	True	1	1	

	class_First	class_Second	class_Third
0	0.0	0.0	1.0
2	0.0	0.0	1.0
3	1.0	0.0	0.0
4	0.0	0.0	1.0
5	0.0	0.0	1.0
..
886	0.0	1.0	0.0
887	1.0	0.0	0.0
888	0.0	0.0	1.0
889	1.0	0.0	0.0
890	0.0	0.0	1.0

[718 rows x 17 columns]

Step 8: Feature Scaling

```
from sklearn.preprocessing import StandardScaler

# Initialize the scaler
scaler = StandardScaler()

# Scale Age and Fare columns
df_titanic[['Age', 'Fare']] = scaler.fit_transform(df_titanic[['age',
'fare']])

df_titanic.head()
```

	survived	pclass	sex	age	fare	embarked	class	who
adult_male \								
0	0	3	male	22.0	7.2500	S	Third	man
True								
2	1	3	female	26.0	7.9250	S	Third	woman
False								
3	1	1	female	35.0	53.1000	S	First	woman
False								
4	0	3	male	35.0	8.0500	S	Third	man
True								
5	0	3	male	28.0	8.4583	Q	Third	man
True								

	embark_town	alive	alone	sex_encoded	who_encoded	class_First	\
0	Southampton	no	False	1	1	0.0	
2	Southampton	yes	True	0	2	0.0	
3	Southampton	yes	False	0	2	1.0	
4	Southampton	no	True	1	1	0.0	
5	Queenstown	no	True	1	1	0.0	

	class_Second	class_Third	Age	Fare
0	0.0	1.0	-0.607611	-0.751265
2	0.0	1.0	-0.207827	-0.700265
3	0.0	0.0	0.691688	2.712961
4	0.0	1.0	0.691688	-0.690821
5	0.0	1.0	-0.007934	-0.659971